

## ▼ Drive Mount and Dataset Unzip

### HEALTH CARE

Predicting pneumonia from chest X-ray images

Dataset link: <https://www.kaggle.com/tolgadincer/labeled-chest-xray-images>

Google Drive link: <https://drive.google.com/file/d/1xYceBz1JM5D4TDNQM02yMDYbe4oHZTlt/view?usp=sharing>

```
1 from google.colab import drive
2 drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)

1 # Unzip updated dataset from drive, with labelled test data
2 !mkdir dataset
3 %cd dataset
4 !unzip "/content/drive/MyDrive/archive.zip"
5 %cd ..
```

## ▼ Common Cells for both Models

**CONSTANTS, Helper Functions, Imports, Optimizer, Data Generator, Class Weights, Test Data Load**

### CONSTANTS

```
1 input_shape = (150,150,3)
2 target_size = (150,150)
3 epochs = 50
4 batch_size = 64
5 patience = 3
6 classes = ('normal', 'pneumonia')
7 train_dir = './dataset/chest_xray/train'
8 test_dir = './dataset/chest_xray/test'
```

### Helper Functions

```
1 def plot_model_accuracy(H):
2     plt.plot(H.history['accuracy'])
3     plt.plot(H.history['val_accuracy'])
4     plt.title('Model Accuracy')
5     plt.ylabel('Accuracy')
6     plt.xlabel('Epoch')
7     plt.legend(['Train Accuracy', 'Validation Accuracy'], loc='upper left', bbox_to_anchor=(1,1))
8     plt.show()
9
10 def plot_model_loss(H):
11     plt.plot(H.history['loss'])
12     plt.plot(H.history['val_loss'])
13     plt.title('Model Loss')
14     plt.ylabel('Loss')
15     plt.xlabel('Epoch')
16     plt.legend(['Train Loss', 'Validation Loss'], loc='upper left', bbox_to_anchor=(1,1))
17     plt.show()
18
19 def plot_model_lr(H):
20     N = np.arange(0, len(H.epoch))
21     plt.style.use('ggplot')
22     plt.figure()
23     plt.plot(N, H.history['accuracy'], label='train_accuracy')
24     plt.plot(N, H.history['val_accuracy'], label='val_accuracy')
25     plt.plot(N, H.history['loss'], label='train_loss')
26     plt.plot(N, H.history['val_loss'], label='val_loss')
27     plt.title('Training loss and accuracy')
28     plt.xlabel('Epoch #')
29     plt.ylabel('Loss/Accuracy')
30     plt.legend()
31     plt.show()
32
33 # ...
```



```

5                                     batch_size = batch_size)
6
7 valid_dataset = valid_datagen.flow_from_directory(directory = train_dir,
8                                     target_size = target_size,
9                                     class_mode = 'categorical',
10                                    subset = 'validation',
11                                    batch_size = batch_size)
12
13 test_dataset = test_datagen.flow_from_directory(directory = test_dir,
14                                     target_size = target_size,
15                                     class_mode = 'categorical',
16                                     batch_size = batch_size)

```

```

Found 4187 images belonging to 2 classes.
Found 1045 images belonging to 2 classes.
Found 624 images belonging to 2 classes.

```

## Evaluate Class Weights (To handle data imbalance)

```

1 train_class_weights = class_weight.compute_class_weight(
2     'balanced',
3     np.unique(train_dataset.classes),
4     train_dataset.classes)
5
6 class_weights = { 0: train_class_weights[0], 1: train_class_weights[1] }
7 print('Class weights (imblanced classes):', class_weights)

Class weights (imblanced classes): {0: 1.938425925925926, 1: 0.6738010943031864}

```

## Load Test Data

```

1 print('Loading test images:')
2 X_test, y_test = load_data(test_dir)

3%|██████████| 6/234 [00:00<00:04, 50.42it/s]Loading test images:
100%|██████████| 234/234 [00:05<00:00, 45.10it/s]
100%|██████████| 390/390 [00:02<00:00, 146.21it/s]

```

## Test Labels to One Hot Encoding

```

1 lb = LabelBinarizer()
2 y_test = lb.fit_transform(y_test)
3 y_test = np.hstack((1 - y_test, y_test))

```

# ▼ Resnet50 (Base Model)

## Resnet50 imagenet

```

1 from keras.applications.resnet50 import ResNet50
2 base_model = ResNet50(input_shape=input_shape, weights='imagenet', include_top=False)

```

## Freezing Layers

```

1 for layer in base_model.layers:
2     layer.trainable=False

```

## Defining Layers

```

1 model_base=Sequential()
2 model_base.add(base_model)
3 model_base.add(GlobalAveragePooling2D())
4 model_base.add(Dense(64, activation='relu'))
5 model_base.add(Dropout(0.3))
6 model_base.add(Dense(2,activation='softmax'))

```

## Model Summary

```

1 model_base.summary()

```

```
1 model_base.summary()
```

Model: "sequential\_12"

Layer (type)	Output Shape	Param #
resnet50 (Functional)	(None, 5, 5, 2048)	23587712
global_average_pooling2d_132	(None, 2048)	0
dense_278 (Dense)	(None, 64)	131136
dropout_19 (Dropout)	(None, 64)	0
dense_279 (Dense)	(None, 2)	130

Total params: 23,718,978  
Trainable params: 131,266  
Non-trainable params: 23,587,712

Model Compile with SGD Optimzer

```
1 model_base.compile(loss='categorical_crossentropy',
2                     optimizer=optimizer,
3                     metrics=['accuracy'])
```

Defining Callbacks

```
1 filepath = './best_base_weights.hdf5'
2
3 early_stopping = EarlyStopping(monitor = 'val_accuracy',
4                                 mode = 'max',
5                                 patience = patience,
6                                 verbose = 1)
7
8 checkpoint      = ModelCheckpoint(filepath,
9                                    monitor = 'val_accuracy',
10                                    mode='max',
11                                    save_best_only=True,
12                                    verbose = 1)
13
14 callback_list = [early_stopping, checkpoint]
```

Model Fitting

```
1 H_base = model_base.fit(train_dataset,
2                           validation_data=valid_dataset,
3                           epochs = epochs,
4                           callbacks = callback_list,
5                           class_weight = class_weights,
6                           verbose = 1)
```

Epoch 1/50  
66/66 [=====] - 53s 772ms/step - loss: 0.7648 - accuracy: 0.4847 - val\_loss: 0.6941 - val\_accuracy: 0.

Epoch 00001: val\_accuracy improved from -inf to 0.43254, saving model to ./best\_base\_weights.hdf5  
Epoch 2/50  
66/66 [=====] - 50s 752ms/step - loss: 0.6924 - accuracy: 0.4798 - val\_loss: 0.6980 - val\_accuracy: 0.

Epoch 00002: val\_accuracy did not improve from 0.43254  
Epoch 3/50  
66/66 [=====] - 49s 749ms/step - loss: 0.6926 - accuracy: 0.4645 - val\_loss: 0.6770 - val\_accuracy: 0.

Epoch 00003: val\_accuracy improved from 0.43254 to 0.74258, saving model to ./best\_base\_weights.hdf5  
Epoch 4/50  
66/66 [=====] - 50s 751ms/step - loss: 0.6884 - accuracy: 0.5729 - val\_loss: 0.6878 - val\_accuracy: 0.

Epoch 00004: val\_accuracy improved from 0.74258 to 0.75694, saving model to ./best\_base\_weights.hdf5  
Epoch 5/50  
66/66 [=====] - 50s 755ms/step - loss: 0.6869 - accuracy: 0.5367 - val\_loss: 0.6707 - val\_accuracy: 0.

Epoch 00005: val\_accuracy did not improve from 0.75694  
Epoch 6/50  
66/66 [=====] - 49s 749ms/step - loss: 0.6851 - accuracy: 0.5937 - val\_loss: 0.7069 - val\_accuracy: 0.

Epoch 00006: val\_accuracy did not improve from 0.75694  
Epoch 7/50  
66/66 [=====] - 49s 751ms/step - loss: 0.6788 - accuracy: 0.5087 - val\_loss: 0.6742 - val\_accuracy: 0.

```
Epoch 00007: val_accuracy improved from 0.75694 to 0.76364, saving model to ./best_base_weights.hdf5
Epoch 8/50
66/66 [=====] - 49s 748ms/step - loss: 0.6872 - accuracy: 0.6037 - val_loss: 0.6725 - val_accuracy: 0.

Epoch 00008: val_accuracy improved from 0.76364 to 0.77703, saving model to ./best_base_weights.hdf5
Epoch 9/50
66/66 [=====] - 50s 751ms/step - loss: 0.6851 - accuracy: 0.5900 - val_loss: 0.6675 - val_accuracy: 0.

Epoch 00009: val_accuracy did not improve from 0.77703
Epoch 10/50
66/66 [=====] - 50s 753ms/step - loss: 0.6829 - accuracy: 0.5870 - val_loss: 0.6581 - val_accuracy: 0.

Epoch 00010: val_accuracy did not improve from 0.77703
Epoch 11/50
66/66 [=====] - 50s 753ms/step - loss: 0.6865 - accuracy: 0.5698 - val_loss: 0.6794 - val_accuracy: 0.

Epoch 00011: val_accuracy did not improve from 0.77703
Epoch 00011: early stopping
```



Load Base Model Weights (if required)

```
1 model_base.load_weights("best_base_weights.hdf5")
```

▼ Resnet50 (Summary)

Test Loss/Accuracy, Training/Validation Graphs, Confusion Matrix, Classification Report & Visual Results


Evaluating Loss and AUC - Test Data

```
1 score = model_base.evaluate(test_dataset)
2 print('Test Loss: ', score[0])
3 print('Test Accuracy: ', score[1])

10/10 [=====] - 6s 588ms/step - loss: 0.6780 - accuracy: 0.7244
Test Loss:  0.6780446767807007
Test Accuracy:  0.7243589758872986
```

Summarize Model Loss

```
1 plot_model_loss(H_base)
```



Summarize Model Accuracy


```
1 plot_model_accuracy(H_base)
```



## Summarize Learning Curve (Accuracy and Loss)



```
1 plot_model_lr(H_base)
```



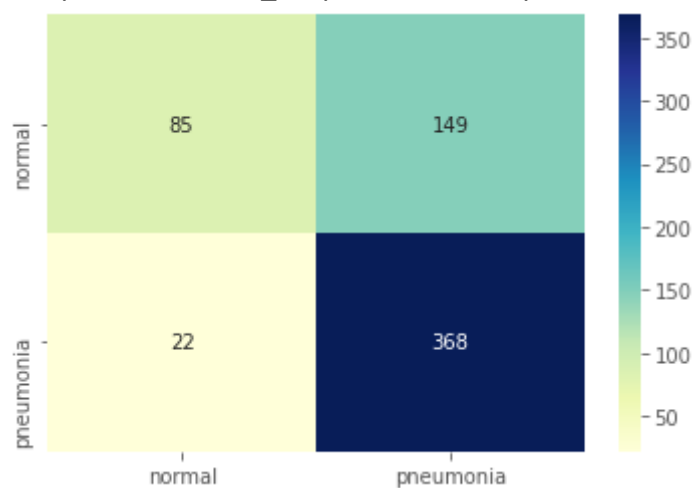
## CONFUSION MATRIX

```
1 # Making prediction
2 y_pred = model_base.predict(X_test)
3 y_true = np.argmax(y_test, axis=-1)
4
5 # Plotting the confusion matrix
6 from sklearn.metrics import confusion_matrix
7 confusion_mtx = confusion_matrix(y_true, np.argmax(y_pred, axis=1))
8 confusion_mtx
```

```
array([[ 85, 149],
       [ 22, 368]])
```

```
1 import seaborn as sns
2 sns.heatmap(confusion_mtx, xticklabels=classes, yticklabels=classes, annot=True, fmt='d', cmap="YlGnBu")
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f7b8fd17898>
```



## Classification Report (Precision, Recall, F1-score, Support)

```
1 from sklearn.metrics import classification_report
2 predictions = model_base.predict(X_test, batch_size=32)
3 print(classification_report(y_test.argmax(axis=1), predictions.argmax(axis=1), target_names=classes))
```


	precision	recall	f1-score	support
normal	0.79	0.36	0.50	234
pneumonia	0.71	0.94	0.81	390
accuracy			0.73	624
macro avg	0.75	0.65	0.66	624
weighted avg	0.74	0.73	0.69	624

## Visual Results

```

1 fig, ax = plt.subplots(nrows=5, ncols=5, sharex=True, sharey=True, figsize=(12,12))
2 num=0
3 for i in range(5):
4     for j in range(5):
5         img = X_test[num]
6         ax[i][j].imshow(img)
7         ohe = one_hot_encoding(y_pred[num])
8         ax[i][j].set_title(get_title(classes[ohe.index(1)], y_test[num]))
9         num += 18
10
11 ax[0][0].set_yticks([])
12 ax[0][0].set_xticks([])
13 plt.tight_layout()
14 plt.show()

```



## ▼ Resnet50 + CBAM (Enhanced Model)

```

1 from keras.layers import GlobalAveragePooling2D, GlobalMaxPooling2D, Reshape, Dense, multiply, Permute, Concatenate, Conv2D, Add,
2 from keras import backend as K
3 from keras.activations import sigmoid
4
5 def attach_attention_module(net, attention_module):
6     net = cbam_block(net)
7     return net
8
9 def cbam_block(cbam_feature, ratio=8):
10     cbam_feature = channel_attention(cbam_feature, ratio)
11     cbam_feature = spatial_attention(cbam_feature)
12     return cbam_feature
13
14

```

```

15 def channel_attention(input_feature, ratio=8):
16     channel_axis = 1 if K.image_data_format() == "channels_first" else -1
17     channel = input_feature.shape[channel_axis]
18
19     shared_layer_one = Dense(channel//ratio,
20                             activation='relu',
21                             kernel_initializer='he_normal',
22                             use_bias=True,
23                             bias_initializer='zeros')
24     shared_layer_two = Dense(channel,
25                             kernel_initializer='he_normal',
26                             use_bias=True,
27                             bias_initializer='zeros')
28
29     avg_pool = GlobalAveragePooling2D()(input_feature)
30     avg_pool = Reshape((1,1,channel))(avg_pool)
31     assert avg_pool.shape[1:] == (1,1,channel)
32     avg_pool = shared_layer_one(avg_pool)
33     assert avg_pool.shape[1:] == (1,1,channel//ratio)
34     avg_pool = shared_layer_two(avg_pool)
35     assert avg_pool.shape[1:] == (1,1,channel)
36
37     max_pool = GlobalMaxPooling2D()(input_feature)
38     max_pool = Reshape((1,1,channel))(max_pool)
39     assert max_pool.shape[1:] == (1,1,channel)
40     max_pool = shared_layer_one(max_pool)
41     assert max_pool.shape[1:] == (1,1,channel//ratio)
42     max_pool = shared_layer_two(max_pool)
43     assert max_pool.shape[1:] == (1,1,channel)
44
45     cbam_feature = Add()([avg_pool,max_pool])
46     cbam_feature = Activation('sigmoid')(cbam_feature)
47
48     if K.image_data_format() == "channels_first":
49         cbam_feature = Permute((3, 1, 2))(cbam_feature)
50
51     return multiply([input_feature, cbam_feature])
52
53 def spatial_attention(input_feature):
54     kernel_size = 7
55
56     if K.image_data_format() == "channels_first":
57         channel = input_feature.shape[1]
58         cbam_feature = Permute((2,3,1))(input_feature)
59     else:
60         channel = input_feature.shape[-1]
61         cbam_feature = input_feature
62
63     avg_pool = Lambda(lambda x: K.mean(x, axis=3, keepdims=True))(cbam_feature)
64     assert avg_pool.shape[-1] == 1
65     max_pool = Lambda(lambda x: K.max(x, axis=3, keepdims=True))(cbam_feature)
66     assert max_pool.shape[-1] == 1
67     concat = Concatenate(axis=3)([avg_pool, max_pool])
68     assert concat.shape[-1] == 2
69     cbam_feature = Conv2D(filters = 1,
70                           kernel_size=kernel_size,
71                           strides=1,
72                           padding='same',
73                           activation='sigmoid',
74                           kernel_initializer='he_normal',
75                           use_bias=False)(concat)
76     assert cbam_feature.shape[-1] == 1
77
78     if K.image_data_format() == "channels_first":
79         cbam_feature = Permute((3, 1, 2))(cbam_feature)
80
81     return multiply([input_feature, cbam_feature])

```

```

1 import keras
2 from keras.layers import Dense, Conv2D, BatchNormalization, Activation
3 from keras.layers import AveragePooling2D, Input, Flatten
4 from keras.optimizers import Adam
5 from keras.callbacks import ModelCheckpoint, LearningRateScheduler
6 from keras.callbacks import ReduceLROnPlateau
7 from keras.preprocessing.image import ImageDataGenerator
8 from keras.regularizers import l2
9 from keras import backend as K
10 from keras.models import Model

```



```

11
12 def resnet_layer(inputs,
13                 num_filters=16,
14                 kernel_size=3,
15                 strides=1,
16                 activation='relu',
17                 batch_normalization=True,
18                 conv_first=True):
19
20     conv = Conv2D(num_filters,
21                 kernel_size=kernel_size,
22                 strides=strides,
23                 padding='same',
24                 kernel_initializer='he_normal',
25                 kernel_regularizer=l2(1e-4))
26
27     x = inputs
28     if conv_first:
29         x = conv(x)
30         if batch_normalization:
31             x = BatchNormalization()(x)
32         if activation is not None:
33             x = Activation(activation)(x)
34     else:
35         if batch_normalization:
36             x = BatchNormalization()(x)
37         if activation is not None:
38             x = Activation(activation)(x)
39         x = conv(x)
40     return x
41
42 def resnet_v1(input_shape, depth, attention_module='cbam_block'):
43     if (depth - 2) % 6 != 0:
44         raise ValueError('depth should be 6n+2 (eg 20, 32, 44 in [a])')
45     # Start model definition.
46     num_filters = 16
47     num_res_blocks = int((depth - 2) / 6)
48
49     inputs = Input(shape=input_shape)
50     x = resnet_layer(inputs=inputs)
51     # Instantiate the stack of residual units
52     for stack in range(3):
53         for res_block in range(num_res_blocks):
54             strides = 1
55             if stack > 0 and res_block == 0: # first layer but not first stack
56                 strides = 2 # downsample
57             y = resnet_layer(inputs=x,
58                             num_filters=num_filters,
59                             strides=strides)
60             y = resnet_layer(inputs=y,
61                             num_filters=num_filters,
62                             activation=None)
63             if stack > 0 and res_block == 0: # first layer but not first stack
64                 # linear projection residual shortcut connection to match
65                 # changed dims
66                 x = resnet_layer(inputs=x,
67                                 num_filters=num_filters,
68                                 kernel_size=1,
69                                 strides=strides,
70                                 activation=None,
71                                 batch_normalization=False)
72
73             y = attach_attention_module(y, attention_module)
74             x = keras.layers.add([x, y])
75             x = Activation('relu')(x)
76             num_filters *= 2
77
78     # Add classifier on top.
79     # v1 does not use BN after last shortcut connection-ReLU
80     x = AveragePooling2D(pool_size=8)(x)
81     x = Flatten()(x)
82     x = Dense(1024, activation='relu')(x)
83     x = BatchNormalization()(x)
84     x = Dense(1024, activation='relu')(x)
85     x = BatchNormalization()(x)
86     x = Dropout(0.5)(x)
87     outputs = Dense(2, activation='softmax')(x)
88
89     # Instantiate model.
90     model = Model(inputs=inputs, outputs=outputs)

```

Resnet 50 + CBAM (Enhanced Model)

```
1 # For ResNet, specify the depth (e.g. ResNet50: depth=50)
2 depth = 50
3 model_enhanced = resnet_v1(input_shape=input_shape, depth=depth, attention_module='cbam_block')
```

Model Summary

```
1 model_enhanced.summary()
```

Model: "model\_14"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_29 (InputLayer)	[(None, 150, 150, 3)]	0	
conv2d_1050 (Conv2D)	(None, 150, 150, 16)	448	input_29[0][0]
batch_normalization_727 (Batch Normalization)	(None, 150, 150, 16)	64	conv2d_1050[0][0]
activation_1041 (Activation)	(None, 150, 150, 16)	0	batch_normalization_727[0][0]
conv2d_1051 (Conv2D)	(None, 150, 150, 16)	2320	activation_1041[0][0]
batch_normalization_728 (Batch Normalization)	(None, 150, 150, 16)	64	conv2d_1051[0][0]
activation_1042 (Activation)	(None, 150, 150, 16)	0	batch_normalization_728[0][0]
conv2d_1052 (Conv2D)	(None, 150, 150, 16)	2320	activation_1042[0][0]
batch_normalization_729 (Batch Normalization)	(None, 150, 150, 16)	64	conv2d_1052[0][0]
global_average_pooling2d_349 (Global Average Pooling)	(None, 16)	0	batch_normalization_729[0][0]
global_max_pooling2d_336 (Global Max Pooling)	(None, 16)	0	batch_normalization_729[0][0]
reshape_672 (Reshape)	(None, 1, 1, 16)	0	global_average_pooling2d_349[0][0]
reshape_673 (Reshape)	(None, 1, 1, 16)	0	global_max_pooling2d_336[0][0]
dense_738 (Dense)	(None, 1, 1, 2)	34	reshape_672[0][0] reshape_673[0][0]
dense_739 (Dense)	(None, 1, 1, 16)	48	dense_738[0][0] dense_738[1][0]
add_672 (Add)	(None, 1, 1, 16)	0	dense_739[0][0] dense_739[1][0]
activation_1043 (Activation)	(None, 1, 1, 16)	0	add_672[0][0]
multiply_672 (Multiply)	(None, 150, 150, 16)	0	batch_normalization_729[0][0] activation_1043[0][0]
lambda_672 (Lambda)	(None, 150, 150, 1)	0	multiply_672[0][0]
lambda_673 (Lambda)	(None, 150, 150, 1)	0	multiply_672[0][0]
concatenate_336 (Concatenate)	(None, 150, 150, 2)	0	lambda_672[0][0] lambda_673[0][0]
conv2d_1053 (Conv2D)	(None, 150, 150, 1)	98	concatenate_336[0][0]
multiply_673 (Multiply)	(None, 150, 150, 16)	0	multiply_672[0][0] conv2d_1053[0][0]
add_673 (Add)	(None, 150, 150, 16)	0	activation_1041[0][0] multiply_673[0][0]
activation_1044 (Activation)	(None, 150, 150, 16)	0	add_673[0][0]

Model Compile with SGD Optimizer

```
1 model_enhanced.compile(loss='binary_crossentropy',
2                         metrics=[tensorflow.keras.metrics.AUC(name = 'accuracy')],
3                         optimizer=optimizer)
```

Data Augmentation (Training Data)

```

1 train_augmented = ImageDataGenerator(rescale = 1.0 / 255.0,
2                                     featurewise_center=False,
3                                     samplewise_center=False,
4                                     featurewise_std_normalization=False,
5                                     samplewise_std_normalization=False,
6                                     zca_whitening=False,
7                                     rotation_range=20,
8                                     width_shift_range=0.2,
9                                     height_shift_range=0.2,
10                                    shear_range=0.2,
11                                    zoom_range=0.2,
12                                    channel_shift_range=0.2,
13                                    fill_mode='nearest',
14                                    horizontal_flip=False,
15                                    validation_split = 0.2)

```

## Flow form Directory (Training & Validation Data)

```

1 train_augmented_dataset = train_augmented.flow_from_directory(directory = train_dir,
2                                                                target_size = target_size,
3                                                                class_mode = 'categorical',
4                                                                subset = 'training',
5                                                                batch_size = batch_size)

```

Found 4187 images belonging to 2 classes.

## Defining Callbacks

```

1 filepath = './best_enhanced_weights.hdf5'
2
3 early_stopping = EarlyStopping(monitor = 'val_accuracy',
4                                mode = 'max',
5                                patience = patience,
6                                verbose = 1)
7
8 checkpoint      = ModelCheckpoint(filepath,
9                                   monitor = 'val_accuracy',
10                                  mode='max',
11                                  save_best_only=True,
12                                  verbose = 1)
13
14 callback_list = [early_stopping, checkpoint]

1 # Fit the model on the batches generated by datagen.flow().
2 H_enhanced = model_enhanced.fit(train_augmented_dataset,
3                                 validation_data=valid_dataset,
4                                 class_weight = class_weights,
5                                 epochs=epochs,
6                                 verbose=1,
7                                 callbacks=callback_list)

```

```

Epoch 1/50
66/66 [=====] - 116s 2s/step - loss: 1.0176 - accuracy: 0.7680 - val_loss: 1.0940 - val_accuracy: 0.55

Epoch 00001: val_accuracy improved from -inf to 0.55956, saving model to ./best_enhanced_weights.hdf5
Epoch 2/50
66/66 [=====] - 98s 1s/step - loss: 0.7820 - accuracy: 0.8856 - val_loss: 0.9551 - val_accuracy: 0.799

Epoch 00002: val_accuracy improved from 0.55956 to 0.79963, saving model to ./best_enhanced_weights.hdf5
Epoch 3/50
66/66 [=====] - 99s 1s/step - loss: 0.7573 - accuracy: 0.9036 - val_loss: 1.0167 - val_accuracy: 0.860

Epoch 00003: val_accuracy improved from 0.79963 to 0.86033, saving model to ./best_enhanced_weights.hdf5
Epoch 4/50
66/66 [=====] - 99s 1s/step - loss: 0.7468 - accuracy: 0.9037 - val_loss: 1.0716 - val_accuracy: 0.840

Epoch 00004: val_accuracy did not improve from 0.86033
Epoch 5/50
66/66 [=====] - 99s 1s/step - loss: 0.7034 - accuracy: 0.9251 - val_loss: 0.7838 - val_accuracy: 0.918

Epoch 00005: val_accuracy improved from 0.86033 to 0.91848, saving model to ./best_enhanced_weights.hdf5
Epoch 6/50
66/66 [=====] - 99s 1s/step - loss: 0.7171 - accuracy: 0.9207 - val_loss: 1.0501 - val_accuracy: 0.829

Epoch 00006: val_accuracy did not improve from 0.91848
Epoch 7/50
66/66 [=====] - 100s 2s/step - loss: 0.7023 - accuracy: 0.9223 - val_loss: 0.6844 - val_accuracy: 0.94

```

```
Epoch 00007: val_accuracy improved from 0.91848 to 0.94817, saving model to ./best_enhanced_weights.hdf5
Epoch 8/50
66/66 [=====] - 99s 1s/step - loss: 0.7003 - accuracy: 0.9232 - val_loss: 0.5939 - val_accuracy: 0.970

Epoch 00008: val_accuracy improved from 0.94817 to 0.97035, saving model to ./best_enhanced_weights.hdf5
Epoch 9/50
66/66 [=====] - 100s 1s/step - loss: 0.6764 - accuracy: 0.9389 - val_loss: 0.5968 - val_accuracy: 0.97

Epoch 00009: val_accuracy improved from 0.97035 to 0.97266, saving model to ./best_enhanced_weights.hdf5
Epoch 10/50
66/66 [=====] - 100s 2s/step - loss: 0.6459 - accuracy: 0.9444 - val_loss: 1.4110 - val_accuracy: 0.68

Epoch 00010: val_accuracy did not improve from 0.97266
Epoch 11/50
66/66 [=====] - 100s 1s/step - loss: 0.6592 - accuracy: 0.9463 - val_loss: 2.0835 - val_accuracy: 0.48

Epoch 00011: val_accuracy did not improve from 0.97266
Epoch 12/50
66/66 [=====] - 99s 1s/step - loss: 0.6251 - accuracy: 0.9535 - val_loss: 0.8991 - val_accuracy: 0.885

Epoch 00012: val_accuracy did not improve from 0.97266
Epoch 00012: early stopping
```

Load Enhanced Model Weights (if required)

```
1 model_enhanced.load_weights("best_enhanced_weights.hdf5")
```

Resnet50 + CBAM (Summary)

Test Loss/Accuracy, Training/Validation Graphs, Confusion Matrix, Classification Report & Visual Results

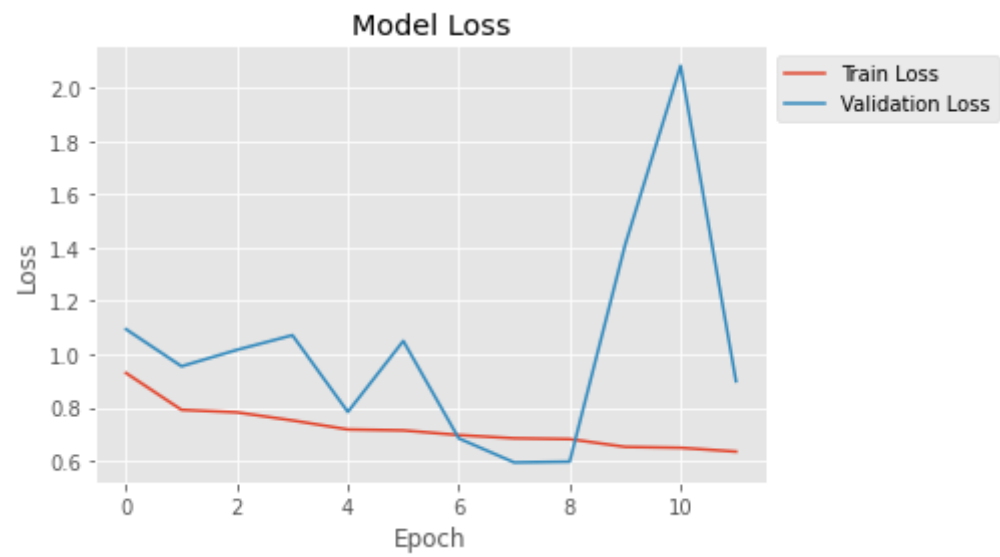
Evaluating Loss and AUC - Test Data

```
1 score = model_enhanced.evaluate(test_dataset)
2 print('Test Loss: ', score[0])
3 print('Test Accuracy: ', score[1])

10/10 [=====] - 7s 633ms/step - loss: 0.6687 - accuracy: 0.9432
Test Loss: 0.6687048673629761
Test Accuracy: 0.9431834816932678
```


Summarize Model Loss

```
1 plot_model_loss(H_enhanced)
```




Summarie Model Accuracy


```
1 plot_model_accuracy(H_enhanced)
```



## Summarize Learning Curve (Accuracy and Loss)



```
1 plot_model_lr(H_enhanced)
```



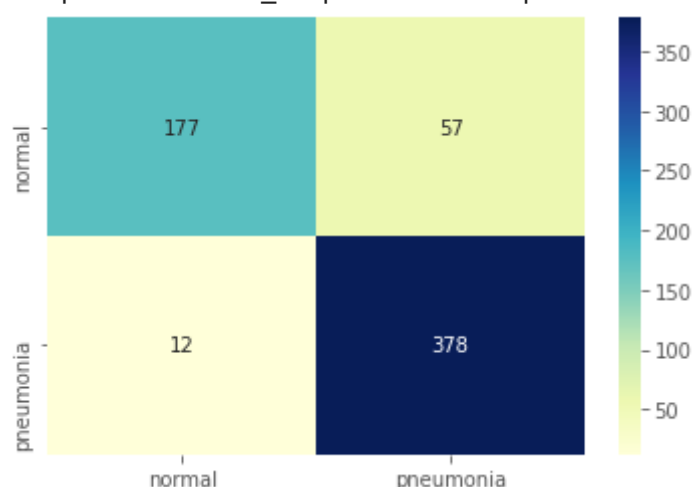
## CONFUSION MATRIX

```
1 # Making prediction
2 y_pred_enhanced = model_enhanced.predict(X_test)
3 y_true = np.argmax(y_test, axis=-1)
4
5 # Plotting the confusion matrix
6 from sklearn.metrics import confusion_matrix
7 confusion_mtx = confusion_matrix(y_true, np.argmax(y_pred_enhanced, axis=1))
8 confusion_mtx

array([[177,  57],
       [ 12, 378]])

1 import seaborn as sns
2 sns.heatmap(confusion_mtx, xticklabels=classes, yticklabels=classes, annot=True, fmt='d', cmap="YlGnBu")
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f79664fdc50>



## Classification Report (Precision, Recall, F1-score, Support)

```
1 from sklearn.metrics import classification_report
2 predictions = model_enhanced.predict(X_test, batch_size=32)
3 print(classification_report(y_test.argmax(axis=1), predictions.argmax(axis=1), target_names=classes))
```

	precision	recall	f1-score	support
normal	0.94	0.76	0.84	234
pneumonia	0.87	0.97	0.92	390
accuracy			0.89	624
macro avg	0.90	0.86	0.88	624
weighted avg	0.89	0.89	0.89	624

## Visual Results

```

1 fig, ax = plt.subplots(nrows=5, ncols=5, sharex=True, sharey=True, figsize=(12,12))
2 num=0
3 for i in range(5):
4     for j in range(5):
5         img = X_test[num]
6         ax[i][j].imshow(img)
7         ohe = one_hot_encoding(y_pred_enhanced[num])
8         ax[i][j].set_title(get_title(classes[ohe.index(1)], y_test[num]))
9         num += 18
10
11 ax[0][0].set_yticks([])
12 ax[0][0].set_xticks([])
13 plt.tight_layout()
14 plt.show()

```

